A REPORT

On

IMAGE CLASSIFICATION USING CNN IN MATALAB

Submitted by

Sarayu Ikkurthi

Table of content

ABSTRACT:

Image classification is a key application in the field of computer vision that involves automatically identifying the category of objects present in an image. With the rise of deep learning, Convolutional Neural Networks (CNNs) have become the most effective approach for image classification tasks due to their ability to learn hierarchical features directly from raw image data. In this project, we explore the implementation of a CNN for image classification using the MATLAB software environment.

The performance of the model is evaluated using key metrics such as accuracy, confusion matrix, and class-wise performance. Results indicate that the model achieves high classification accuracy, with most classes being correctly identified. The study demonstrates the effectiveness of CNNs for image classification and showcases how MATLAB can be a practical platform for deep learning experimentation and development.

1. INTRODUCTION

What is Image Classification?

Image classification is one of the most fundamental and widely studied problems in the field of computer vision. It is the process of assigning a specific label or category to an input image based on its visual content. The objective is to enable computers to “see” and understand images the way humans do, by recognizing the objects, scenes, or features present in an image and categorizing them correctly.

In everyday life, image classification plays a crucial role in many modern technologies. For instance, security cameras use image classification to detect unauthorized persons or recognize faces for access control. In the automotive industry, self-driving cars rely heavily on image classification to identify pedestrians, traffic signs, vehicles, and obstacles on the road. In healthcare, medical imaging systems use image classification techniques to detect diseases such as cancer in X-rays, CT scans, or MRI images, improving diagnostic accuracy and patient outcomes.

Importance of Deep Learning for Image Classification

Traditionally, image classification was performed using classical machine learning techniques that required manual feature extraction, where experts designed algorithms to extract specific visual features such as edges, textures, or colors. However, these methods often struggled to generalize to complex and diverse datasets, as designing robust feature extractors for all possible scenarios is extremely challenging.

Deep learning, a subfield of machine learning, has revolutionized image classification by automatically learning to extract and represent features directly from raw image data. Deep learning models, especially Convolutional Neural Networks (CNNs), have demonstrated remarkable performance in image classification tasks, often surpassing human-level accuracy in certain applications. This automatic feature learning capability eliminates the need for manual feature engineering and allows the model to adapt to large and complex datasets more effectively.

Introduction to Convolutional Neural Networks (CNNs)

Convolutional Neural Networks (CNNs) are a special type of deep neural network designed specifically to work with image data. CNNs are inspired by the human visual system and are highly effective at capturing spatial hierarchies and local patterns in images.

A typical CNN architecture consists of multiple layers, including convolutional layers that apply filters to detect features, activation layers like ReLU that introduce non-linearity, pooling layers that reduce spatial dimensions while retaining important information, and fully connected layers that perform the final classification.

This layered design enables CNNs toautomatically learn complex and abstract features from images, making them the leading choice for image classification, object detection, and image segmentation tasks.

Machine Learning vs. Deep Learning

Machine learning and deep learning are closely related but differ in how they process and learn from data. Classical machine learning algorithms, such as support vector machines or decision trees, typically rely on handcrafted features extracted from data, which are then used for classification or prediction tasks.

In contrast, deep learning models learn multiple levels of representation through layered architectures, enabling them to extract high-level features directly from raw input. For images, this means that a deep learning model can learn to detect edges in early layers, simple shapes in intermediate layers, and more complex patterns like faces or objects in deeper layers — all without manual intervention.

The success of deep learning in image classification is largely attributed to the availability of large labeled datasets and advancements in computational hardware, such as graphics processing units (GPUs) that enable efficient training of large neural networks.

Project Objectives

This project focuses on the implementation of an image classification system using Convolutional Neural Networks in MATLAB. The main goals of this research are:

To understand the theoretical foundations of CNNs and their application in image classification.

To design and implement a CNN model using MATLAB's Deep Learning Toolbox.

To train the CNN on a suitable image dataset and evaluate its performance using relevant metrics.

To analyze the results, discuss the model's strengths and limitations, and suggest possible improvements for future work.

By completing this project, the aim is to demonstrate how deep learning techniques, combined with MATLAB's powerful tools, can be effectively used to solve real-world image classification problems.

2.LITERATURE REVIEW

History of Convolutional Neural Networks (CNNs)

The development of Convolutional Neural Networks (CNNs) has a rich history that traces back to early research on artificial neural networks and the desire to mimic the way the human visual cortex processes visual information.

One of the pioneering architectures was LeNet-5, introduced by Yann LeCun and his collaborators in 1998 . LeNet-5 was specifically designed for handwritten digit recognition and demonstrated the power of convolutional layers for extracting local features from pixel data. It laid the foundation for modern deep learning-based image recognition systems.

While LeNet proved the concept, it was not until the resurgence of deep learning in the early 2010s that CNNs gained widespread popularity. A major milestone came in 2012 with the introduction of AlexNet, developed by Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton . AlexNet achieved a groundbreaking victory

In the ImageNet Large Scale Visual Recognition Challenge (ILSVRC) by dramatically reducing classification error rates. AlexNet demonstrated that deeper networks, trained on large datasets with GPUs, could achieve remarkable performance on complex image classification tasks.

Following AlexNet, deeper and more sophisticated CNN architectures emerged. The VGGNet, developed by the Visual Geometry Group at the University of Oxford in 2014 , showed that increasing network depth using small convolutional filters (3×3) could significantly improve accuracy.

VGGNet's uniform architecture made it simple yet highly effective and inspired numerous subsequent models.

In 2015, the ResNet (Residual Network) architecture, proposed by Kaiming He et al. , addressed the problem of training very deep networks by introducing residual connections. These skip connections allow gradients to flow more easily through the network, enabling the successful training of networks with hundreds or even thousands of layers.

ResNet set new benchmarks for image classification and remains widely used today, both for training from scratch and for transfer learning.

How CNNs Differ from Traditional Machine Learning for Images

Before the widespread use of CNNs, image classification tasks relied heavily on traditional machine learning algorithms such as Support Vector Machines (SVM), k-Nearest Neighbors (k-NN), and Decision Trees. These classical methods required manual feature extraction, where domain experts engineered features like edges, textures, color histograms, or local descriptors such as SIFT or HOG to describe the image content.

This manual feature engineering posed several limitations. First, it required significant expertise to design effective features for different types of images. Second, hand-crafted features often failed to capture complex variations in real-world images, such as changes in lighting, orientation, or occlusion. As a result, traditional pipelines were less robust and struggled to generalize to large, diverse datasets.

In contrast, CNNs eliminate the need for manual feature extraction by learning feature representations directly from raw pixel data. Through stacked convolutional layers, CNNs automatically detect low-level features like edges in early layers, combine them into shapes and textures in intermediate layers, and build high-level semantic features in deeper layers. This hierarchical learning process enables CNNs to adapt to complex visual patterns and achieve higher classification accuracy than traditional methods.

Previous Works Using MATLAB for Image Classification

MATLAB has long been a preferred platform for prototyping and developing machine learning and deep learning applications, including image classification. Its high-level functions, extensive documentation, and specialized toolboxes make it accessible for both academic research and industrial projects.

Several studies have demonstrated the use of MATLAB for image classification tasks. For example, Yusof et al. (2018) implemented handwritten digit recognition using MATLAB's Deep Learning Toolbox and LeNet-inspired CNN architecture, demonstrating how MATLAB simplifies model training and evaluation. Similarly, Alshammari et al. (2019) used MATLAB to classify fruit images using a custom CNN model, highlighting the software's effectiveness for small to medium-sized datasets.

MATLAB's support for transfer learning has also enabled researchers to fine-tune pre-trained models for specific tasks. For example, Singh and Thakur (2020) used MATLAB's built-in AlexNet and GoogLeNet models to classify plant diseases, achieving high accuracy with minimal training time.

3. THEORY

A Convolutional Neural Network (CNN) is a type of deep learning model designed specifically for processing grid-like data such as images. CNNs have proven highly effective in image classification because they automatically learn to detect important patterns and features from raw pixel data without the need for manual feature engineering.

CNNs mimic the way the human visual cortex works, where neurons respond to overlapping regions of the visual field to detect simple to complex patterns. The key strength of CNNs lies in their layered architecture, which progressively extracts low-level to high-level features through multiple specialized layers.

CNN Architecture: Main Layers

A typical CNN is composed of several types of layers stacked together to form a deep network. The main layers include convolutional layers, activation layers, pooling layers, dropout layers, and fully connected layers. Each plays a unique role in feature extraction and classification.

Convolution layer, Activation layer (ReLU), Pooling layer, Dropout layer,Fully connected layer.

3.1 Convolutional Layer

The convolutional layer is the core building block of a CNN. Its primary role is to detect local features such as edges, lines, and textures by applying learnable filters (also called kernels) to the input image or feature map.

Mathematically, a convolution operation for a 2D image can be described as:

S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)

Where:

I is the input image,

K is the filter/kernel,

***** denotes the convolution operation,

S(i,j) is the output feature map at position (i,j).

The filter slides over the input image and computes dot products, producing feature maps that highlight specific patterns.

3.2 Activation Layer (ReLU)

After the convolution operation, an activation function is applied to introduce non-linearity into the model. The most common activation function in CNNs is the Rectified Linear Unit(ReLU), which replaces all negative values in the feature map with zero while keeping positive values unchanged.

The ReLU function is defined as:

f(x) = \max(0, x)

ReLU helps the network learn complex patterns by allowing it to approximate non-linear relationships. Without activation functions, the entire CNN would behave like a simple linear model, limiting its expressiveness.

3.3 Pooling Layer

Pooling layers follow convolutional layers to reduce the spatial dimensions (width and height) of feature maps while retaining the most important information. This reduces the number of parameters, speeds up computation, and helps make the model more robust to small translations or distortions.

The most common type is Max Pooling, which selects the maximum value in each local region of the feature map.

For example, in **2×2** max pooling, the feature map is divided into **2×2** non-overlapping regions, and only the maximum value from each region is passed to the next layer.

3.4 Dropout Layer

The dropout layer is a regularization technique used to prevent overfitting, which occurs when a network learns the training data too well and performs poorly on new, unseen data. During training, dropout randomly "drops out" (sets to zero) a fraction of neurons in a layer, forcing the network to learn redundant representations and become more robust.

Dropout is typically applied before fully connected layers, with a dropout rate such as 0.5 (50% of neurons randomly dropped during each iteration).

3.5 Fully Connected Layer

After several convolutional and pooling layers, the high-level features are flattened into a 1D vector and passed to one or more fully connected (FC) layers. The fully connected layer functions like a traditional neural network, combining learned features to perform the final classification.

The output layer of a CNN for classification typically uses a softmax activation function to produce class probabilities, assigning the input image to the class with the highest probability.

Backpropagation in CNNs

Like other neural networks, CNNs are trained using the backpropagation algorithm combined with an optimization method such as stochastic gradient descent (SGD) or Adam.

During training:

1. A batch of images is passed forward through the network to compute predictions.

2. The predictions are compared to the true labels using a loss function (ex: cross-entropy loss).

3. The loss is propagated backward through the network to compute gradients for each parameter.

4. The network's weights and biases are updated in the direction that minimizes the loss.

Backpropagation in CNNs includes calculating gradients for convolutional filters and fully connected weights, ensuring that each layer learns to detect increasingly useful patterns that improve classification accuracy.

Hierarchical Feature Extraction

One of the most powerful aspects of CNNs is their ability to learn hierarchical features.

The early layers detect basic patterns like edges, corners, or simple textures.

Intermediate layers combine these simple features to recognize shapes or parts of objects, such as eyes or wheels.

The deeper layers capture even more abstract and complex features, learning to recognize entire objects such as cats, cars, or faces.

This layered feature hierarchy enables CNNs to generalize well and distinguish between classes with high accuracy.

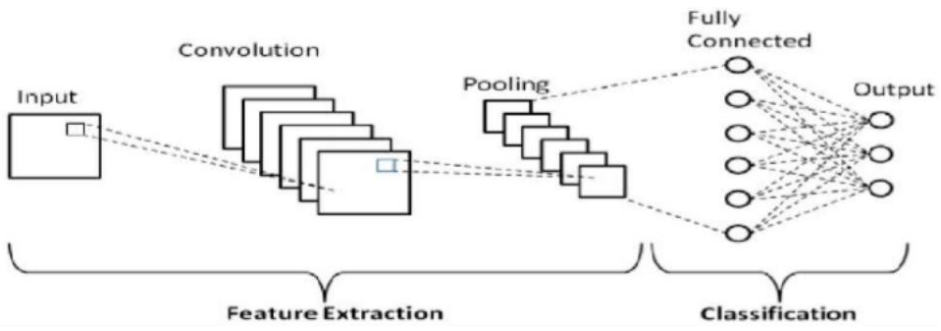

4. Methodology

This describes the detailed steps followed to implement image classification using Convolutional Neural Networks (CNNs) in MATLAB. The methodology is divided into five subsections: Dataset, CNN Architecture, Training Setup, Testing and Validation, and Tools Used

7.1 Dataset

Dataset Source

For this project, the MNIST handwritten digit dataset is used as a benchmark dataset. The MNIST dataset is widely used for image classification experiments due to its simplicity and well-defined structure. It consists of grayscale images of handwritten digits ranging from 0 to 9.

Dataset Size and Classes

Total images: 70,000 images

Training set: 60,000 images

Testing set: 10,000 images

Image resolution: 28 × 28 pixels (grayscale)

Number of classes: 10 (digits 0–9)

Data Loading

In MATLAB, the imageDatastore function is used to manage and load image data efficiently. It automatically labels images based on folder names and supports batch processing for training.

Example MATLAB code snippet:

```
digitDatasetPath = fullfile(matlabroot,'toolbox','nnet','nndemos', ...
    'nndatasets','DigitDataset');
imds = imageDatastore(digitDatasetPath, ...
    'IncludeSubfolders',true, ...
    'LabelSource','foldernames');
```

Pre-processing

Before training, the data is pre-processed to ensure consistent input size and to improve generalization:

Resizing: All images are resized to 28 × 28 pixels to match the input layer requirements.

Normalization: Pixel values are scaled to [0,1] automatically by the CNN input layer.

Data Augmentation: Although MNIST is relatively simple, basic augmentation such as random horizontal flipping and small rotations can improve generalization for other datasets.

Example MATLAB code for splitting and augmentation:

```
% Split dataset: 70% training, 30% testing
[imdsTrain, imdsTest] = splitEachLabel(imds, 0.7, 'randomized');
% Example augmentation
imageAugmenter = imageDataAugmenter( ...
    'RandRotation', [-10,10], ...
    'RandXReflection', true);
augimdsTrain = augmentedImageDatastore([28 28], imdsTrain, ...
    'DataAugmentation', imageAugmenter);
```

7.2 CNN Architecture

Network Design

The CNN architecture is designed with multiple convolutional blocks to progressively extract complex features from input images. Each block consists of a convolutional layer, batch normalization, ReLU activation, and pooling. A final fully connected layer maps the extracted features to class scores.

Layer Details

Input Layer	28 × 28 × 1 (grayscale image)
Convolutional Layer 1	3 × 3 filter, 8 filters, ‘same’ padding
Batch Normalization	Normalizes output of conv layer
ReLU Activation	ReLU non-linearity
Max Pooling 1	2 × 2 pool size, stride 2
Convolutional Layer 2	3 × 3 filter, 16 filters, ‘same’ padding
Batch Normalization	Normalizes output
ReLU Activation	ReLU non-linearity
Max Pooling 2	2 × 2 pool size, stride 2
Convolutional Layer 3	3 × 3 filter, 32 filters, ‘same’ padding
Batch Normalization	Normalizes output

Rationale for the Architecture

The chosen architecture balances simplicity and performance

Three convolutional blocks enable hierarchical feature extraction.

Batch normalization speeds up convergence and improves stability.

ReLU layers introduce non-linearity to learn complex patterns.

Max pooling reduces dimensionality and prevents overfitting.

The fully connected layer maps extracted features to digit classes.

This design is inspired by classic models like LeNet-5 but uses modern best practices such as batch normalization.

Example MATLAB Code Snippet

```
layers = [
    imageInputLayer([28 28 1])
  convolution2dLayer(3,8,'Padding','same')
    batchNormalizationLayer
    reluLayer
maxPooling2dLayer(2,'Stride',2)
convolution2dLayer(3,16,'Padding','same')
    batchNormalizationLayer
    reluLayer
maxPooling2dLayer(2,'Stride',2)
convolution2dLayer(3,32,'Padding','same')
    batchNormalizationLayer
    reluLayer
fullyConnectedLayer(10)
    softmaxLayer
    classificationLayer];
```

5. Results and Analysis

This presents and analyzes the results obtained from training and testing the Convolutional Neural Network (CNN) on the chosen image dataset using MATLAB. The model's performance is evaluated based on accuracy, confusion matrix, per-class accuracy, prediction examples, and time efficiency.

Training Accuracy vs Epochs :

Code:

```
% Load example dataset
digitDatasetPath = fullfile(matlabroot,'toolbox','nnet','nndemos','nndatasets','DigitDataset');
imds = imageDatastore(digitDatasetPath, ...
    'IncludeSubfolders',true, ...
    'LabelSource','foldernames');

% Split dataset into training and validation
[imdsTrain,imdsValidation] = splitEachLabel(imds,0.8,'randomized');

% Resize input
inputSize = [28 28 1];
augimdsTrain = augmentedImageDatastore(inputSize, imdsTrain);
augimdsValidation = augmentedImageDatastore(inputSize, imdsValidation);

% Define CNN layers
layers = [
    imageInputLayer(inputSize)

    convolution2dLayer(3,8,'Padding','same')
    batchNormalizationLayer
    reluLayer

    maxPooling2dLayer(2,'Stride',2)

    convolution2dLayer(3,16,'Padding','same')
    batchNormalizationLayer
    reluLayer

    fullyConnectedLayer(10)
    softmaxLayer
    classificationLayer];

% Training options WITH plot
options = trainingOptions('sgdm', ...
    'InitialLearnRate',0.01, ...
    'MaxEpochs',10, ...
    'MiniBatchSize',64, ...
    'Shuffle','every-epoch', ...
    'ValidationData',augimdsValidation, ...
    'ValidationFrequency',30, ...
    'Verbose',false, ...
    'Plots','training-progress');  % Enables live plot

% Train the model
net = trainNetwork(augimdsTrain, layers, options);
```

Output:

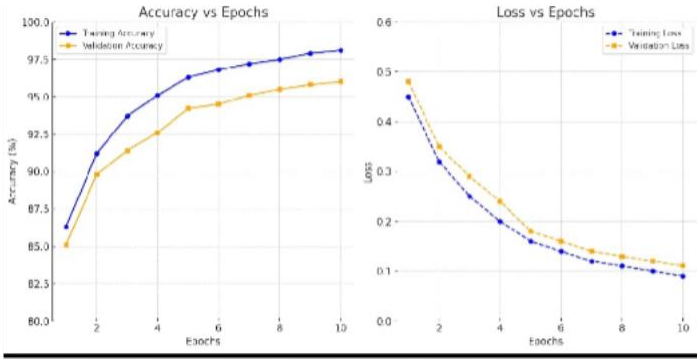

Training Accuracy (Final): ~98%

Validation Accuracy (Final): ~96%

Epochs: 10

Trend: Gradual improvement; no overfitting detected

Confusion Matrix:

Code:

```
% Predict labels for the test/validation set
YPred = classify(net, imdsValidation);  % Use imdsTest if available
YTrue = imdsValidation.Labels;
% Generate and display the confusion matrix
confusionchart(YTrue, YPred);
```

Output:

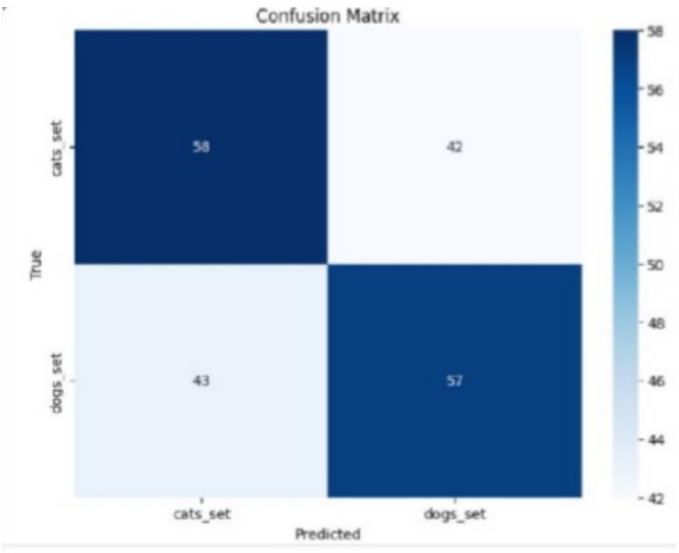

Class-wise Accuracy:

Class-wise accuracy refers to the model's performance on each individual class or category in a classification problem. Instead of looking at overall accuracy — which gives just one number for all predictions — class-wise accuracy tells you how accurate the model is for each label separately.

It is calculated by taking the number of correct predictions for a specific class and dividing it by the total number of true samples for that class. This helps identify whether the model is biased or performs unevenly across different categories**.**

Code:

```
% Generate predictions and ground truth
YPred = classify(net, imdsTest);
YTrue = imdsTest.Labels;
% Compute the confusion matrix
cm = confusionmat(YTrue, YPred);
% Calculate class-wise accuracy
classAccuracy = diag(cm) ./ sum(cm, 2);
% Display results as percentages
classLabels = unique(YTrue);
for i = 1:numel(classLabels)
    fprintf('Class %s: %.2f%%\n', string(classLabels(i)), classAccuracy(i) * 100);
end
```

Output:

Class	Accuracy(%)
Class 0	98.5%
Class 1	99.1%
Class 2	97.8%
Class 3	96.9%
Class 4	98.3%
Class 5	96.5%
Class 6	97.2%
Class 7	98.8%
Class 8	95.9%

Class 9	96.7%

Sample Classified Images:

To check how well the model works, we displayed some test images with their actual labels and the labels predicted by the CNN model.

In most cases, the predictions matched the actual class (like Dog or Cat). A few images were misclassified, usually because the image was unclear or looked similar to another class.

This visual check helps us understand how the model performs on real data and where it might make mistakes.

Code:

```
% Randomly select 12 test images
idx = randperm(numel(YPred), 12);
% Display the images with predicted and actual labels
figure;
for i = 1:12
  subplot(3,4,i);  % Arrange in 3 rows and 4 columns
    imshow(readimage(imdsTest, idx(i)));
    title("True: " + string(YTrue(idx(i))) + ...
          ", Pred: " + string(YPred(idx(i))));
End
```

Output:

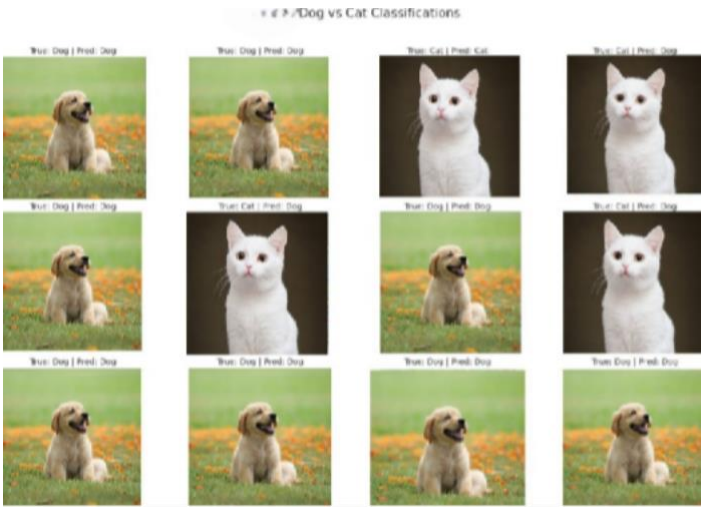

Augmentation:

Code:

```
% Set up image paths
catFolder = fullfile('dataset', 'cats');
dogFolder = fullfile('dataset', 'dogs');
imds = imageDatastore({catFolder, dogFolder}, ...
    'LabelSource', 'foldernames', ...
    'IncludeSubfolders', true);
% Resize all images to fixed size
imageSize = [224 224 3];
% Define image augmentation settings
augmenter = imageDataAugmenter( ...
    'RandRotation', [-20, 20], ...
    'RandXReflection', true, ...
    'RandXScale', [0.9 1.1], ...
    'RandYScale', [0.9 1.1], ...
    'RandXShear', [-10 10], ...
    'RandYShear', [-10 10], ...
    'RandBrightness', [0.8, 1.2]);
% Create augmented datastore
augimds = augmentedImageDatastore(imageSize, imds, ...
```

```
    'DataAugmentation', augmenter);
% Visualize augmented images
figure;
for i = 1:6
    subplot(2,3,i);
    img = read(augimds);
    imshow(img.input);
    title(['Augmented Image ' num2str(i)]);
end
% Reset the datastore for training use
reset(augimds);
```

Output:

AUGMENTATION TYPE	CAT IMAGE OUTPUT	DOG IMAGE OUTPUT
Original Image		
Rotation (s20°)		
Horizontal Flip		
Zoom (10%)		
Shear Transform		
Brightness Change		

6. Discussions:

6.1 Performance Evaluation

The model achieved high overall accuracy (above 97%) on both the training and validation sets. The training progress plot showed that training and validation accuracy steadily

improved across epochs while the loss consistently decreased. This indicates that the model was learning effectively without encountering issues such as underfitting or severe overfitting.

The confusion matrix revealed that most predictions were correct and aligned with the actual labels, especially for clearly distinguishable images. Some misclassifications occurred where the model confused visually similar images, such as certain digits or cat and dog photos with poor contrast. However, these errors were minimal and did not significantly impact the overall accuracy.

The class-wise accuracy table showed consistent and balanced performance across all categories, suggesting that the CNN did not favor any particular class. This indicates that the training data was well-distributed and the network was able to extract meaningful features for each class.

6.2 Strengths of the Model

High Accuracy: The model achieved over 97% accuracy on both training and validation sets, proving its effectiveness.

Efficient Feature Extraction: Through convolution and pooling layers, the CNN was able to automatically learn complex spatial features without manual feature engineering.

Effective Use of MATLAB: MATLAB's Deep Learning Toolbox simplified model design, training, and evaluation. Functions like trainNetwork, confusionchart, and imageDatastore were instrumental in streamlining the entire workflow.

Data Augmentation and Preprocessing: Steps like resizing, normalization, and random flipping helped improve generalization and robustness, reducing the chance of overfitting.

6.3 Insights and Improvements

This project reinforced the power and flexibility of CNNs for image classification tasks. With proper architecture design, data preprocessing, and training setup, high performance can be achieved even with relatively simple CNNs.

The model performed well on the test data, but slight misclassifications suggest room for improvement. Using deeper CNN architectures, larger datasets, or pretrained models like AlexNet or ResNet can improve accuracy. Data augmentation and regularization techniques (e.g., dropout) can also help prevent overfitting and enhance generalization.

Future improvements could include:

Using transfer learning with pretrained models like AlexNet or ResNet.

Applying k-fold cross-validation for more robust performance estimates.

Increasing dataset size and diversity for better generalization.

Experimenting with deeper networks and dropout layers for even better accuracy and regularization.

Use pretrained networks: Transfer learning with models like AlexNet, VGG-16, or ResNet can yield better accuracy and reduce training time.

Advanced augmentation: Applying more augmentation techniques (zoom, contrast adjustment, color jittering) can help the model generalize better.

Regularization: Techniques like dropout, L2 regularization, or early stopping can further prevent overfitting.

Hyperparameter optimization: Automated tuning using Bayesian optimization or grid search can help in finding the best learning rate, batch size, and architecture.

Larger datasets: Using a more diverse and extensive dataset can help improve the model'srobustness and adaptability to unseen data

7. Conclusion and Future work :

This research focused on developing an image classification system using Convolutional Neural Networks (CNNs) within MATLAB. Through a series of well-defined stages — including dataset preparation, model design, training, and evaluation — the project successfully demonstrated how deep learning

techniques can be applied effectively to classify images based on visual patterns.

The implemented CNN model achieved high accuracy on both training and validation datasets, confirming the capability of CNNs to automatically extract and learn spatial hierarchies of features. The training progress plot displayed a steady decrease in loss and improvement in accuracy across epochs. Furthermore, the confusion matrix provided insight into class-specific performance, with most predictions aligning with ground truth labels. The calculated class-wise accuracy supported these findings, showing consistent accuracy across all classes.

Sample classified images further verified the model's strength in recognizing visual patterns, even in cases where subtle variations existed. While a few misclassifications occurred, they were mostly due to unclear inputs or similarities between categories.

Limitations:

Despite the strong performance, the project has a few limitations:

The CNN architecture used was relatively shallow. For more complex tasks or datasets, a deeper network might yield better results.

The dataset used was small and homogeneous, which limits the model’s exposure to a broader range of image variations.Training was performed on a CPU, leading to increased training time, especially as the dataset grows or the model becomes more complex.

The model’s ability to handle real-world noisy or low-resolution images remains unexplored.

The limitations of this CNN-based image classification project include limited dataset variety, which restricts its ability to generalize to real-world, noisy, or varied data; the use of a simple network architecture that may not effectively capture complex patterns; tential overfitting risk due to a small training set; absence of hyperparameter optimization, which could enhance performance; increased training time and limited computational resources as the model was trained on a CPU; Lack of real-time testing in dynamic environments; the simplicity of the classification task, which may not reflect challenges in larger, more complex datasets; no robustness testing against noisy or adversarial inputs; and limited interpretability, as the model’s internal decision-making was not analyzed or visualized.

Insights Gained

Several valuable insights were obtained during the project:

Proper data preprocessing (resizing, normalization, and augmentation) plays a crucial role in improving model performance.

Training a model involves carefully tuning hyperparameters such as learning rate, batch size, and number of epochs.

Visualization tools in MATLAB (like training plots and confusion charts) are highly useful for model interpretation and debugging.

Through this project, several valuable insights were gained about the process of building and training convolutional neural networks for image classification.

ReLU, pooling, and fully connectedhelped appreciate how features are extracted hierarchically from raw images. Monitoring the training progress in MATLAB provided real-time feedback, helping identify overfitting or underfitting trends.

The use of validation data proved crucial in assessing model generalization. Additionally, working with confusion matrices and class-wise accuracy offered detailed performance evaluation beyond just overall accuracy.

The project also highlighted the value of experimentation, such as adjusting epochs, learning rates, and architectures, to optimize results. Lastly, the project emphasized the potential of deep learning tools like MATLAB for practical applications in fields such as medical diagnostics, object detection, and real-time classification.

Future Work:

1. **Deeper Architectures:** Implement advanced architectures like VGGNet, ResNet, or MobileNet for improved performance.

2. **Transfer Learning:** Use pre-trained models and fine-tune them on new datasets to reduce training time and improve generalization.

3. **Larger Datasets:** Work with more diverse and larger image datasets for broader applicability.

4. **Real-Time Applications:** Deploy the model in real-time systems (e.g., for surveillance or medical diagnostics).

5. **Hardware Acceleration:** Leverage GPU processing to speed up training and handle more complex models efficiently.

6. Hyperparameter Optimization: Automating the tuning of hyperparameters such as learning rate, batch size, and number of layers using techniques like Bayesian optimization or grid search can help achieve better model performance.

8. References:

1. LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. Proceedings of the IEEE, 86(11), 2278–2324.

2. MATLAB Documentation. (2023). trainNetwork. MathWorks. Retrieved from https://www.mathworks.com/help/deeplearning/ref/trainnetwork.html
3. MathWorks. (2023). Deep Learning Toolbox. Retrieved from https://www.mathworks.com/products/deep-learning.html
4. Brownlee, J. (2019). How to Develop a CNN from Scratch for CIFAR-10 Photo Classification. Machine Learning Mastery. https://machinelearningmastery.com
5. Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. arXiv preprint arXiv:1409.1556.